

# POLS 4724: Replication Experiment

Aman Choudhri

December 23, 2024

In this paper, I replicate (and test for robustness to modeling choices) the results of Dymnicki et al. The original paper is a 46 school randomized controlled trial of a holistic educational intervention known as Safe Communities Safe Schools (SCSS). The SCSS program was initially developed in 1999 at the University of Colorado Boulder’s Center for the Study and Prevention of Violence (CSPV) following the Columbine shooting, with the goal of addressing mental/behavioral health concerns and facilitating prosocial behavior. See *Safe Communities Safe Schools - CSPV* for more information.

The variant of the program evaluated in Dymnicki et al. consisted of three main components: a team sited within each school dedicated to implementing the comprehensive program over multiple years; “capacity building” to help school staff take advantage of data in their decision-making; and the implementation of an evidence-based “action plan” to develop a school’s support systems for students.

According to the authors, this particular SCSS variant had previously not yet been evaluated in the literature. They approach their experiment with the goals of studying how well the SCSS program is implemented, as measured by qualitative rubric assessments from annual CSPV staff visits, and the effects of the program on school climate, student behavior, and academic outcomes.

## 1 Research Design

The trial was conducted on  $N = 46$  middle schools in Colorado, with a student population of  $S = 62590$  students. The study employed a staggered cohort design with two cohorts ( $N_1 = 10, N_2 = 36$ ) beginning one year apart. Within each cohort, schools were matched into pairs based on school characteristics and test scores, with one school from each pair randomly assigned to either the treatment or waitlist control condition.

The evaluation followed schools over a three-year period (2016-2018), with control schools scheduled to implement the treatment two years after their cohort’s treatment group began implementation. Due to the staggered nature of the cohorts, the analysis focused on program effectiveness during the first two years of implementation for each cohort to maintain comparability.

The key outcome variables of interest are:

- *School climate*, with subcategories like “peer norms” and “violence indicators”. Climate was measured through annual student and staff surveys.
- *Student behavior*, as measured by class attendance, truancy, and suspension rates. These data were obtained as school-level averages from the Colorado Department of Education.
- *Student achievement*, as measured by math and reading standardized test scores. Again, the results were provided directly by the Colorado Department of Education.

## 2 Replication and Robustness

In this replication, we focus on studying outcome variables that are “downstream” of school climate—the more tangible phenomena of attendance, truancy, suspensions, and academic achievement. Significant improvements on school climate, which we might think of as a latent, amorphous construct, can reasonably be expected to manifest in terms of these phenomena. So this restricted focus may still shed light on the main social outcomes of interest to the SCSS program.

Unfortunately, baseline data for suspensions were not accessible through the public data page, so the authors analyses of the impact of the program on that outcome variable cannot be reproduced. So for behavioral outcomes, we analyze attendance and truancy only.

### 2.1 Randomization Check

Before proceeding with the replication, we first perform a randomization check to assess whether treatment was effectively randomized within cohorts. The randomization check is performed using by regressing the treatment indicator on baseline school covariates (like the percentage of students eligible for free or reduced price lunch) and the cohort indicator. See Table A for the full set of school-level covariates examined, along with summary statistics for each.

The  $F$ -test fails to reject the null hypothesis of no correlation between covariate profiles and treatment with  $F = 0.531, p = 0.856$ . As a result, we can be reasonably confident that the randomization procedure was performed correctly and that there are no strange inaccuracies from the matching process.

### 2.2 School-Level Outcomes

First, we match the paper and estimate the treatment effect on attendance and truancy. The authors estimate two separate average treatment effects per outcome variable: one after one year of treatment, and one after two years of treatment. Following the paper’s authors, we estimate each ATE by pooling across the cohort groups and using a regression estimator:

$$Y_s^{(t)} \sim \epsilon_s + \beta_0^{(t)} + \beta_1^{(t)} Z_s + \beta_2^{(t)} y_s^{(0)} + X_s^T \beta_3 + \sum_{j=2}^{23} \gamma_j I_{s \in \text{pair } i}.$$

In this notation,  $Y_s^{(t)}$  is the outcome variable for school  $s$  in year  $t$ ;  $Z_s$  is of course the treatment indicator; the variable  $y_s^{(0)}$  represents the *baseline* measurement of the outcome in the year prior to treatment;  $X_s$  is a matrix of the school-level covariates discussed in the previous section; and each  $I_{s,i}$  is an indicator variable representing whether school  $s$  belongs to pair  $i$ .

All regression results accord with the analyses from the authors, finding no treatment effects across the board. The coefficients, standard errors, and p-values for the attendance and truancy outcome variables across years are displayed in Table 1.

An important methodological question is of missing data: one school closed down during the experiment, leading to one complete missing row. Another school became “nonresponsive” in the second year of the study. The authors removed both schools and their pairs; if a significant treatment effect had emerged, it may have been interesting to handle the attrition more carefully, especially for the nonresponsive school. Since the findings are null, however, we proceed.

Both attendance and truancy having similarly null treatment effect findings does make sense, since they largely capture the same phenomenon. It’s interesting that the authors originally chose to treat them as two separate outcomes, as Figure 1 shows quite a strong correlation between the two.

Table 1: School-Level Outcome Treatment Effects

Outcome	Model	Year 1			Year 2		
		$\beta^{(1)}$	SE	$p$	$\beta^{(2)}$	SE	$p$
Attendance	Original	0.010	0.0050	0.277	0.000	0.0050	0.949
	Replication	0.006	0.0052	0.277	0.002	0.0050	0.697
Truancy	Original	0.000	0.0040	0.261	0.000	0.0040	0.439
	Replication	-0.004	0.0038	0.261	-0.005	0.0044	0.305

Note: SE = standard error. Original results from Dymnicki et al.

Correlation between Attendance and Truancy Outcomes




Figure 1: Relationship between Attendance and Truancy, by Treatment Year

## 2.3 Individual-Level Outcomes

Next, we replicate the analysis of individual-level data on student achievement. Achievement, in this context, is measured by standardized outcome measures on state-wide reading and math examinations. The fact that the outcome measures are standardized state-wide potentially poses an interesting spillover effect violation. For a large enough sample size, it cannot be the case that students in all schools have constant positive treatment effects. There are  $N_{\text{population}} = 289$  middle schools in Colorado, per the Department of Education, meaning this experiment obtained outcome measurements from  $N/N_{\text{population}} = 15.9\%$  of all schools. This is a nontrivial fraction, so it is worth keeping this fact about the outcome measurements in mind in the following analysis.

Again, the authors estimate two separate treatment effects per outcome measure, one for a one-year program treatment and one for a two-year treatment. Again, a regression model is used, with baseline achievement scores used as individual-level covariates.

The authors also include individual-level demographic information as covariates, many of which have been redacted in the public dataset. We estimate a model using only the available covariates (Male/Female, White/Non-White) to understand the sensitivity of the treatment effect estimate to the model specification.

The results from the regressions are shown in Table 2. In general, the results are similar. Effect sizes are generally near to the original paper’s. However, the obtained standard errors under these modeling choices are *far* higher in general. In fact, the standard errors can be so large as to render findings that were significant under the original paper insignificant.

This pattern is potentially due to the lack of individual-level covariates: the original model’s additional demographic covariates may have helped explain variation in test scores, leading to tighter estimates. By contrast, our model may be running into an identification issue with too many degrees of freedom relative to the effective sample size. Without individual covariates to help explain variability within schools, the model might be struggling to separate different sources of variation.

Table 2: Effects on Reading and Math Academic Achievement Test Scores

Outcome	Grade	Model	Year 1			Year 2		
			$\beta$	SE	$p$	$\beta$	SE	$p$
Reading	7	Original	0.10	0.021	0.000	–	–	–
		Replication	0.176	0.0560	0.00274	–	–	–
	8	Original	0.12	0.023	0.000	0.11	0.025	0.000
		Replication	0.142	0.0723	0.0697	0.092	0.070	0.2231
Math	7	Original	0.03	0.020	0.076	–	–	–
		Replication	0.032	0.0343	0.371	–	–	–
	8	Original	-0.01	0.027	0.739	0.01	0.024	0.795
		Replication	-0.007	0.0050	0.8887	0.00129	0.0573	0.9823

Note: SE = standard error. Original results from Dymnicki et al.

## 2.4 Robustness Check: Other Modeling Approaches

To attempt to repair the precision problem with the model we specified, we now attempt a different model specification. This has the two-fold benefit of 1) probing whether the large standard errors

can be repaired, and 2) understanding the extent to which estimated treatment effects are robust to other sensible modeling choices.

Rather than fitting two separate models for each grade level to estimate one-year treatment effects, we now try pooling data across grades while allowing for grade-specific treatment effects through interaction terms. This approach increases statistical power by sharing information across grades while still allowing for different treatment effects.

The results are displayed in Table 3. We observe notably precise treatment effects than our previous specifications. For reading scores, we find large positive effects in both grades: about 0.21 points in Grade 7 ( $p < 0.001$ ) and 0.15 in Grade 8 ( $p = 0.002$ ). These estimates are not only more precise than our previous replication attempts but also suggest larger effects than the original paper reported.

The math results show an interesting pattern of effect heterogeneity: while Grade 7 shows no effect (in accordance with the previous replication and original results), Grade 8 shows a modest positive effect of 0.065 standard deviations ( $p = 0.024$ ). This grade-level heterogeneity in math impacts was masked in the previous specifications, highlighting the value of this modeling approach in uncovering nuanced treatment effects.

Table 3: One-Year Treatment Effects by Grade Level (Interaction Model)

Outcome	Grade	$\beta$	SE	$p$ -value	Sample	
					$N$	% Treated
Reading	7	0.215	0.046	<0.001	3,789	63.9%
	8	0.150	0.047	0.002	3,513	64.9%
Math	7	-0.012	0.025	0.636	3,789	63.9%
	8	0.065	0.026	0.024	3,513	64.9%

Note: Results from pooled model with treatment-by-grade interactions.

### 3 Conclusion

The replication effort yielded mixed but instructive results. For school-level outcomes like attendance and truancy, our analysis closely matched the original findings, confirming null effects across both treatment years. However, our attempt to replicate the student achievement analysis revealed the crucial role of modeling choices and data availability in treatment effect estimation.

The original achievement analysis relied on individual-level demographic covariates that were redacted in the public dataset. Our initial replication attempt, using only available covariates, produced similar point estimates but with substantially larger standard errors, making some previously significant effects statistically insignificant. This highlighted how omitted individual-level controls, while not necessarily biasing treatment effect estimates, can reduce precision and blur the lines of statistical significance.

However, adopting a different modeling approach that pooled information across grades while allowing for heterogeneous treatment effects yielded notably different results. This specification not only improved precision but also uncovered potentially larger—but also potentially implausible—treatment effects in reading. The contrast between these results and both the original paper and our first replication attempt underscores just how important methodological choices can be when analyzing even data from a randomized experiment.

## References

- Dymnicki, Allison B., et al. “Assessing Implementation and Effects Associated with a Comprehensive Framework Designed to Reduce School Violence: A Randomized Controlled Trial”. *Journal of School Violence* 20, no. 4 (Oct. 2021): 458–470. ISSN: 1538-8220, 1538-8239, visited on 12/23/2024. <https://doi.org/10.1080/15388220.2021.1952078>.
- Safe Communities Safe Schools - CSPV*. <https://cspv.colorado.edu/what-we-do/safeschools/>. Visited on 12/23/2024.

## A School Covariate Summary Statistics

Table 4: School-Level Covariate Summary Statistics

	Mean	Std. Dev.
Proportion Eligible for Free/Reduced Lunch	0.596	0.277
Student-Teacher Ratio	16.008	3.303
% Black	7.318	9.763
% Hispanic	47.564	28.692
% White	37.901	29.858
Small 6th and 7th Grades (less than 200 students)	0.304	0.465
Large 6th and 7th Grades (more than 700 students)	0.196	0.401
Also an Elementary School	0.261	0.444
Also a High School	0.087	0.285